Candidate Shortlist Report

Role	Senior Machine Learning Engineer
Domain	FinTech / AI
Seniority	Senior
Generated	14 May 2026, 14:44
Candidates Screened	6
Model Used	gpt-4o (2024-08-06)

Executive Summary — Ranked Shortlist

#	Name	Score /10	Semantic Sim	Recommendation
1	Arjun Sharma	8.57	0.823	STRONG HIRE
2	Priya Nair	7.78	0.712	HIRE
3	Ananya Das	6.12	0.534	CONSIDER
4	Rahul Verma	5.25	0.398	CONSIDER
5	Deepak Mehta	3.25	0.187	NO HIRE
6	Sneha Kulkarni	2.58	0.098	NO HIRE

Candidate Detail Scorecards

#1 — Arjun Sharma Total: 8.57 / 10 STRONG HIRE

Exceptional candidate with 5 years of directly relevant ML experience in FinTech. Strong technical depth across the full stack — from model training to production deployment. AWS-certified with publications. Top pick for this role.

Dimension	Weight	Score	Visual	Justification
Skills Match	30%	9.0		Covers Python, PyTorch, TensorFlow, LangChain, Docker, K8s — 90%+ skills match.
Experience Relevance	25%	8.5		5 years in ML at FinTech companies; directly relevant domain and seniority.
Education & Certs	15%	9.0		M.Tech from IIT; AWS ML Specialty certification — exceeds requirements.
Project / Portfolio	15%	8.0		Strong portfolio: fraud detection system, RAG pipeline, open-source contribution.
Communication Quality	15%	8.0		Resume is well-structured with clear impact metrics and action verbs.
WEIGHTED TOTAL		8.57		

Skill Gaps: GCP Vertex AI

#2 — Priya Nair Total: 7.78 / 10 HIRE

Strong ML engineer with 4 years of experience and solid technical skills. Domain is Healthcare rather than FinTech, but skills are highly transferable. Would ramp up quickly.

Dimension	Weight	Score	Visual	Justification
Skills Match	30%	8.5		Strong Python, PyTorch, and LangChain skills. Missing Kubernetes experience.
Experience Relevance	25%	7.5		4 years ML experience but in Healthcare domain, not FinTech.
Education & Certs	15%	8.0		M.Tech in Data Science; GCP certification — meets requirements.
Project / Portfolio	15%	7.0		Good portfolio with NLP projects, but no FinTech-specific work.
Communication Quality	15%	7.5		Clear and structured resume, good use of quantitative achievements.
WEIGHTED TOTAL		7.78		

Skill Gaps: Kubernetes · FinTech domain · MLflow

#3 — Ananya Das Total: 6.12 / 10 CONSIDER

Junior candidate with potential but falls short on experience (2.5 years vs 4 required) and lacks several required skills. Could be a fit for a mid-level role.

Dimension	Weight	Score	Visual	Justification
Skills Match	30%	6.0		Has Python and TensorFlow but lacks LangChain, RAG, and MLOps skills.
Experience Relevance	25%	5.5		2.5 years experience — below the 4-year minimum. Adjacent domain.
Education & Certs	15%	7.0		B.Tech in CS from a reputed university. No certifications yet.
Project / Portfolio	15%	6.5		2 academic ML projects. No production or industry portfolio.
Communication Quality	15%	6.0		Resume is functional but lacks impact metrics and specificity.
WEIGHTED TOTAL		**6.12**		

Skill Gaps: LangChain · RAG · Docker · Kubernetes · MLflow · LLM fine-tuning

---

#4 — Rahul Verma Total: 5.25 / 10 CONSIDER

Data analytics background with some ML exposure. Significant skill gaps in deep learning, LLMs, and MLOps. Better suited for a data analyst or junior ML role.

Dimension	Weight	Score	Visual	Justification
Skills Match	30%	5.5		Knows Python and basic ML but lacks deep learning, LLM, and MLOps skills.
Experience Relevance	25%	5.0		3 years in data analytics, not ML engineering. Different role type.
Education & Certs	15%	6.0		B.Tech in IT — meets minimum. No relevant certifications.
Project / Portfolio	15%	4.5		Dashboard and BI projects — not ML model development.
Communication Quality	15%	5.5		Adequate resume but lacks technical depth in descriptions.
WEIGHTED TOTAL		**5.25**		

Skill Gaps: PyTorch · TensorFlow · LangChain · LLM fine-tuning · Docker · Kubernetes · MLflow · RAG

#5 — Deepak Mehta Total: 3.25 / 10 NO HIRE

Experienced frontend developer but completely misaligned with the ML Engineer role. No relevant ML skills, experience, or portfolio.

Dimension	Weight	Score	Visual	Justification
Skills Match	30%	3.0		Frontend developer — Python is secondary. No ML/DL/LLM skills.
Experience Relevance	25%	2.0		4 years of web development — completely different domain and role.

Education & Certs	15%	5.0		B.Tech in CS meets minimum but no ML-related coursework.
Project / Portfolio	15%	2.5		React/Node.js projects — no ML or data science portfolio.
Communication Quality	15%	6.0		Well-written resume but for a completely different role.
WEIGHTED TOTAL		3.25		

Skill Gaps: Python (advanced) · PyTorch · TensorFlow · LangChain · RAG · LLM · MLOps · Docker · Kubernetes

---

#6 — Sneha Kulkarni

Total: 2.58 / 10 NO HIRE

Strong marketing professional but a complete mismatch for an ML Engineer role. No technical skills, no CS education, no relevant projects.

Dimension	Weight	Score	Visual	Justification
Skills Match	30%	2.5		Marketing background — no technical ML skills listed.
Experience Relevance	25%	1.5		3 years in digital marketing — zero ML engineering overlap.
Education & Certs	15%	3.0		MBA in Marketing — does not meet technical education requirement.
Project / Portfolio	15%	1.5		Marketing campaigns and SEO projects — no ML relevance.
Communication Quality	15%	7.0		Excellent writing quality but wrong domain entirely.
WEIGHTED TOTAL		2.58		

Skill Gaps: Python · PyTorch · TensorFlow · LangChain · RAG · LLM · MLOps · Docker · Kubernetes · SQL

HR Override Log

Candidate ID	Overridden By	Original	Adjusted	Reason	Timestamp
cand_003	Priya Sharma – HR Lead	6.12	6.80	Candidate showed strong culture fit and learning agility in preliminary phone screen.	2026-05-10 14:23:45

Methodology & Transparency Appendix

LLM Model: gpt-4o (2024-08-06)
Agent Framework: LangGraph (multi-agent graph with 4 specialised nodes)
Scoring Mode: Dual — Semantic similarity (SentenceTransformers all-MiniLM-L6-v2 cosine) combined with LLM rubric reasoning in JSON mode.
Rubric Weights: Skills Match 30% · Experience Relevance 25% · Project / Portfolio 15% · Education & Certs 15% · Communication Quality 15%
Hallucination Mitigation: Pydantic-validated structured output at every agent node; LLM operates in JSON mode; dimension scores are cross-checked against embedding similarity.
PII Handling: Email and phone are extracted via regex locally and masked before being sent to the cloud LLM; re-injected post-response.
Human-in-the-Loop: HR may override any score via the UI; all overrides are logged with reason and timestamp in this report.
Disclaimer: This report is a decision-support tool. All final hiring decisions must be reviewed and approved by a qualified HR professional.